

```
In [1]: import warnings
warnings.filterwarnings("ignore")
```

# COURSE PROJECT :

## Predicting Final Grades from Student Data

### Introduction

The objective of this project is to predict the final grade of students using multiple linear regression on a dataset of student achievement in two Portuguese schools. We use the Statsmodels and Patsy modules for this task with Python version  $\geq 3.8$ . The dataset was sourced from the UCI Machine Learning Repository at <http://archive.ics.uci.edu/ml/datasets/Student+Performance> (<http://archive.ics.uci.edu/ml/datasets/Student+Performance>) (FUBUTEC 2008). This report is organized as follows:

- [Overview](#) section describes the dataset used and the features in this dataset.
- [Data Preparation](#) section covers data cleaning and data preparation steps.
- [Data Exploration](#) section explores dataset features and their inter-relationships.
- [Statistical Modeling & Performance Evaluation](#) section first fits a full multiple linear regression model and performs diagnostic checks. Next, we perform backwards variable selection using p-values to obtain a reduced model, after which we perform another set of diagnostic checks on the reduced model.
- [Summary & Conclusions](#) section provides a summary of our work and presents our findings.

### Overview

#### Data Source

Our dataset contains data on the achievement of students in secondary education of two Portuguese schools. The dataset was collected from <http://archive.ics.uci.edu/ml/datasets/Student+Performance> (<http://archive.ics.uci.edu/ml/datasets/Student+Performance>), using `student-por.csv`. Our dataset has 649 instances, with 33 attributes, which can be verified with `df.dtypes`.

```
In [2]: import pandas as pd

df = pd.read_csv('Data.csv', sep=';')
pd.set_option('display.max_columns', None)
```

Our dataset was only one column of data including the column title with only the semi-colon (;) symbol separating values, and so it needed to be separated when being read in order to create the separate columns with their respective row values using

```
sep = ';'


```

This is how our dataset currently looks:

```
In [3]: df.sample(10)
```

```
Out[3]:
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	reaso
596	MS	F	17	U	GT3	T	4	2	other	other	cours
603	MS	F	18	R	LE3	A	4	2	teacher	other	reputatic
571	MS	M	19	U	GT3	T	2	1	at_home	other	cours
517	MS	F	16	R	LE3	T	1	2	other	other	reputatic
238	GP	F	16	U	GT3	T	4	4	teacher	services	horr
38	GP	F	15	R	GT3	T	3	4	services	health	cours
484	MS	F	16	U	LE3	A	2	2	other	other	horr
392	GP	F	17	U	GT3	T	3	2	other	other	horr
455	MS	F	15	U	GT3	T	2	1	at_home	other	horr
291	GP	M	18	U	LE3	T	2	1	services	other	cours

## Project Objective

Our goal is to predict our target feature, `G3`, within an acceptable margin of error using linear regression.

## Target Feature

Our target feature is `G3`, which is a discrete numerical feature, and represents the students final grade for a specific course project (Math or Portuguese).

## Descriptive Features

The variable descriptions below are from the `student_por` file:

- `school` - student's school (binary: 'GP' - Gabriel Pereira or 'MS' - Mousinho da Silveira)
- `sex` - student's sex (binary: 'F' - female or 'M' - male)
- `age` - student's age (numeric: from 15 to 22)
- `address` - student's home address type (binary: 'U' - urban or 'R' - rural)
- `famsize` - family size (binary: 'LE3' - less or equal to 3 or 'GT3' - greater than 3)
- `Pstatus` - parent's cohabitation status (binary: 'T' - living together or 'A' - apart)
- `Medu` - mother's education (numeric: 0 - none, 1 - primary education (4th grade), 2 - 5th to 9th grade, 3 - secondary education or 4 - higher education)
- `Fedu` - father's education (numeric: 0 - none, 1 - primary education (4th grade), 2 - 5th to 9th grade, 3 - secondary education or 4 - higher education)
- `Mjob` - mother's job (nominal: 'teacher', 'health' care related, civil 'services' (e.g. administrative or police), 'at\_home' or 'other')
- `Fjob` - father's job (nominal: 'teacher', 'health' care related, civil 'services' (e.g. administrative or police), 'at\_home' or 'other')
- `reason` - reason to choose this school (nominal: close to 'home', school 'reputation', 'course' preference or 'other')
- `guardian` - student's guardian (nominal: 'mother', 'father' or 'other')
- `traveltime` - home to school travel time (numeric: 1 - <15 min., 2 - 15 to 30 min., 3 - 30 min. to 1 hour, or 4 - >1 hour)
- `studytime` - weekly study time (numeric: 1 - <2 hours, 2 - 2 to 5 hours, 3 - 5 to 10 hours, or 4 - >10 hours)
- `failures` - number of past class failures (numeric: n if  $1 \leq n < 3$ , else 4)
- `schoolsup` - extra educational support (binary: yes or no)
- `famsup` - family educational support (binary: yes or no)
- `paid` - extra paid classes within the course subject (Math or Portuguese) (binary: yes or no)
- `activities` - extra-curricular activities (binary: yes or no)
- `nursery` - attended nursery school (binary: yes or no)
- `higher` - wants to take higher education (binary: yes or no)
- `internet` - Internet access at home (binary: yes or no)
- `romantic` - with a romantic relationship (binary: yes or no)
- `famrel` - quality of family relationships (numeric: from 1 - very bad to 5 - excellent)
- `freetime` - free time after school (numeric: from 1 - very low to 5 - very high)
- `goout` - going out with friends (numeric: from 1 - very low to 5 - very high)
- `Dalc` - workday alcohol consumption (numeric: from 1 - very low to 5 - very high)
- `Walc` - weekend alcohol consumption (numeric: from 1 - very low to 5 - very high)
- `health` - current health status (numeric: from 1 - very bad to 5 - very good)
- `absences` - number of school absences (numeric: from 0 to 93)

**these grades are related with the course subject, Math or Portuguese:**

- `G1` - first period grade (numeric: from 0 to 20)
- `G2` - second period grade (numeric: from 0 to 20)
- `G3` - final grade (numeric: from 0 to 20, output target)

## Feature Set

Inspection of the feature descriptions from the `student - por` file allows the removal of features that represent similar data or deemed to have low predictive power.

For instance, features `freetime` and `goout` are deemed to represent similar data, thus feature `goout` is removed preliminarily.

A description of each feature that we will be using in our dataset is presented below in a table format:

name	datatype	units	description
school	binary	'GP' - Gabriel Pereira or 'MS' - Mousinho da Silveira	Student's school
sex	binary	binary: 'F' - female or 'M' - male	Student's gender
address	binary	'U' - urban or 'R' - rural	Student's address (urban/rural)
Medu	numeric	0 - none, 1 - primary education (4th grade), 2 - 5th to 9th grade, 3 - secondary education or 4 - higher education	Highest education achieved by student's mother
Fedu	numeric	0 - none, 1 - primary education (4th grade), 2 - 5th to 9th grade, 3 - secondary education or 4 - higher education	Highest education achieved by student's father
traveltime	numeric	1 - <15 min., 2 - 15 to 30 min., 3 - 30 min. to 1 hour, or 4 - >1 hour	Time spent travelling to and from school
studytime	numeric	1 - <2 hours, 2 - 2 to 5 hours, 3 - 5 to 10 hours, or 4 - >10 hours	Time spent studying over a week
failures	numeric	n if $1 \leq n < 3$ , else 4	Student's total number of past class failures
schoolsup	binary	yes/no	Extra educational support from the school
famsup	binary	yes/no	Extra educational support from the family
paid	binary	yes/no	Extra tutoring classes for Math/Portuguese
higher	binary	yes/no	Student's intention of higher education post -high school
internet	binary	yes/no	Student's access to internet from home
famrel	numeric	1-5	Quality of Student's relationships with his/her family
freetime	numeric	1-5	Amount of free time student has after school
health	numeric	1-5	Current health status
absences	numeric	0-93	Total number of school absences
G3	numeric	0-20	Final grade

## Data Preparation

## Preliminaries

```
In [4]: # Importing modules
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
import statsmodels.formula.api as smf
import scipy.stats as stats
import patsy
%matplotlib inline
%config InlineBackend.figure_format = 'retina'
plt.style.use("ggplot")
```

## Data Cleaning and Transformation

```
In [5]: print(f"Shape of the dataset is {df.shape} \n")
        print(f"Data types are below where 'object' indicates a string type: ")
        print(df.dtypes)
```

Shape of the dataset is (649, 33)

Data types are below where 'object' indicates a string type:

school	object
sex	object
age	int64
address	object
famsize	object
Pstatus	object
Medu	int64
Fedu	int64
Mjob	object
Fjob	object
reason	object
guardian	object
traveltime	int64
studytime	int64
failures	int64
schoolsup	object
famsup	object
paid	object
activities	object
nursery	object
higher	object
internet	object
romantic	object
famrel	int64
freetime	int64
goout	int64
Dalc	int64
Walc	int64
health	int64
absences	int64
G1	int64
G2	int64
G3	int64
dtype:	object

---

Since our response variable is the final grade (G3), we do not need data on the first two periods so we will remove them:

```
In [6]: df = df.drop(['G1', 'G2'], axis = 1)
```

- The student's age is irrelevant information as the range is too small for there to be any significant impact to the student's final age
- The parent's cohabitation status, parent's job, and the reason for selecting the school the student attends, provides little to no value to our dataset, hence we will remove them.
- The guardian of the student would only be relevant if the parents were apart, and considering that not all student's parents are apart, we will also remove this.
- The columns: freetime, studytime, traveltime, goout, and activities all refer to how time was spent outside of school hours. Since we don't need all the specific details of time spent outside of school hours, only keeping freetime, traveltime and studytime would give us all the information we need.
- Whether a student attended nursery school or not would be relevant to their current grades, especially if they're receiving school support, family support and/or paid extra classes.

```
In [7]: df = df.drop(['age', 'Pstatus', 'Mjob', 'Fjob', 'reason', 'guardian', 'goout', 'activities', 'nursery'], axis = 1)
```

Workday alcohol consumption (Dalc) and weekend alcohol consumption (Walc) can be added together to create weekly alcohol consumption (Wkalc), a rating out of 10 (from 0-very low to 8-very high):

```
In [8]: df['Wkalc'] = df['Walc'] + df['Dalc'] - 2
df = df.drop(columns=['Walc', 'Dalc'])
df['Wkalc'].describe()
```

```
Out[8]: count      649.000000
mean         1.782743
std          1.992411
min           0.000000
25%           0.000000
50%           1.000000
75%           3.000000
max           8.000000
Name: Wkalc, dtype: float64
```

## Fixing numerical variables that don't begin with 0:

Our model would make a lot more sense if our numerical variables (eg: famrel is a rating between 1 and 5) began with 0 instead of 1. If variables began with 0, our equation for our data model would be a lot more simplified and our intercept can then be explained as the expected age without the influence of the other variables.

Such variables are:

- traveltime
- studytime
- famrel
- freetime
- health

```
In [9]: df['traveltime'] = df['traveltime'] - 1
df['studytime'] = df['studytime'] - 1
df['famrel'] = df['famrel'] - 1
df['freetime'] = df['freetime'] - 1
df['health'] = df['health'] - 1
```

## Discretising very large ranged numerical variables:

We will check the variable `absences` using the `value_counts` method in Pandas.

```
In [10]: df['absences'].value_counts().sort_index()
```

```
Out[10]: 0      244
1       12
2     110
3        7
4       93
5       12
6       49
7        3
8       42
9        7
10      21
11        5
12      12
13        1
14        8
15        2
16      10
18        3
21        2
22        2
24        1
26        1
30        1
32        1
Name: absences, dtype: int64
```



Let's save a copy of `df['absences']` to `absences` just in case we need to visualise the unmodified data at another time.

```
In [11]: absences = df['absences'].copy()
```

Since the range is such a large value in the variable `absences`, we would need to discretise the data into bins for it to have more of a significant impact on the model.

```
In [12]: df['absences'] = pd.cut(df['absences'], bins = 5, labels=['very low', 'low', 'medium', 'high', 'very high'])
```

Let's see how our values look now using the `value_counts` method in Pandas.

```
In [13]: df['absences'].value_counts()
```

```
Out[13]: very low    527
         low        90
         medium     24
         high        5
         very high   3
         Name: absences, dtype: int64
```

It seems perfect. Let's now perform integer encoding such that *very low* is 0, *low* is 1, *medium* is 2, *high* is 3 and *very high* is 4.

```
In [14]: level_mapping = {'very low': 0, 'low': 1, 'medium': 2, 'high': 3,
                          'very high': 4}
         df['absences'] = df['absences'].replace(level_mapping)

         df.sample(5)
```

```
Out[14]:
```

	school	sex	address	famsize	Medu	Fedu	traveltime	studytime	failures	schoolsuj
584	MS	F	R	GT3	0	0	1	0	0	ni
128	GP	M	R	GT3	4	4	0	0	0	ni
620	MS	F	U	LE3	4	4	0	1	0	ni
12	GP	M	U	LE3	4	4	0	0	0	ni
303	GP	F	U	GT3	3	3	0	2	0	ni

## Checking for Missing Values

```
In [15]: print(f"\nNumber of missing values for each feature:")
print(df.isnull().sum())
```

Number of missing values for each feature:

```
school      0
sex         0
address     0
famsize     0
Medu       0
Fedu       0
traveltime  0
studytime   0
failures    0
schoolsup   0
famsup      0
paid        0
higher      0
internet    0
romantic    0
famrel      0
freetime    0
health      0
absences    0
G3          0
Wkalc       0
dtype: int64
```

**No missing attributes for any of the features** so no need to remove any rows.

```
In [16]: print(f'Now the number of columns are {df.shape[1]}. The dataset c
          urrently looks like:')
df.head()
```

Now the number of columns are 21. The dataset currently looks like:

Out[16]:

	school	sex	address	famsize	Medu	Fedu	traveltime	studytime	failures	schoolsup
0	GP	F	U	GT3	4	4	1	1	0	yes
1	GP	F	U	GT3	1	1	0	1	0	no
2	GP	F	U	LE3	1	1	0	1	0	yes
3	GP	F	U	GT3	4	2	0	2	0	no
4	GP	F	U	GT3	3	3	0	1	0	no

## Summary Statistics

```
In [17]: from IPython.display import display, HTML
display(HTML('<b>Table 1: Summary of continuous features</b>'))
df.describe(include='int64')
```

**Table 1: Summary of continuous features**

Out[17]:

	Medu	Fedu	traveltime	studytime	failures	famrel	freetime
<b>count</b>	649.000000	649.000000	649.000000	649.000000	649.000000	649.000000	649.000000
<b>mean</b>	2.514638	2.306626	0.568567	0.930663	0.221880	2.930663	2.180277
<b>std</b>	1.134552	1.099931	0.748660	0.829510	0.593235	0.955717	1.051093
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
<b>25%</b>	2.000000	1.000000	0.000000	0.000000	0.000000	3.000000	2.000000
<b>50%</b>	2.000000	2.000000	0.000000	1.000000	0.000000	3.000000	2.000000
<b>75%</b>	4.000000	3.000000	1.000000	1.000000	0.000000	4.000000	3.000000
<b>max</b>	4.000000	4.000000	3.000000	3.000000	3.000000	4.000000	4.000000

```
In [18]: display(HTML('<b>Table 2: Summary of categorical features</b>'))
df.describe(include='object')
```

**Table 2: Summary of categorical features**

Out[18]:

	school	sex	address	famsize	schoolsup	famsup	paid	higher	internet	romanti
<b>count</b>	649	649	649	649	649	649	649	649	649	64
<b>unique</b>	2	2	2	2	2	2	2	2	2	
<b>top</b>	GP	F	U	GT3	no	yes	no	yes	yes	n
<b>freq</b>	423	383	452	457	581	398	610	580	498	41

## Data Exploration

### Numerical features

#### Searching for Outliers

We will check if any of the numerical features have any outliers based on Table 1: Summary of continuous features .

- Medu is expected to contain values between 0 to 4 . Based on Table 1, the minimum is 0 and the maximum is 4 and, hence, contains no outliers .
- Fedu is expected to contain values between 0 to 4 . Based on Table 1, the minimum is 0 and the maximum is 4 and, hence, contains no outliers .
- travelttime is expected to contain values between 0 to 3 . Based on Table 1, the minimum is 0 and the maximum is 3 and, hence, contains no outliers .
- studytime is expected to contain values between 0 to 3 . Based on Table 1, the minimum is 0 and the maximum is 3 and, hence, contains no outliers .
- failures is expected to contain values between 0 to 4 . Based on Table 1, the minimum is 0 and the maximum is 3 and, hence, contains no outliers .
- famrel is expected to contain values between 0 to 4 . Based on Table 1, the minimum is 0 and the maximum is 4 and, hence, contains no outliers .
- freetime is expected to contain values between 0 to 4 . Based on Table 1, the minimum is 0 and the maximum is 4 and, hence, contains no outliers .
- health is expected to contain values between 0 to 4 . Based on Table 1, the minimum is 0 and the maximum is 4 and, hence, contains no outliers .
- absences is expected to contain values between 0 to 4 . Based on Table 1, the minimum is 0 and the maximum is 4 and, hence, contains no outliers .
- G3 is expected to contain values between 0 to 20 . Based on Table 1, the minimum is 0 and the maximum is 19 and, hence, contains no outliers .
- Wkalc is expected to contain values between 0 to 8 . Based on Table 1, the minimum is 0 and the maximum is 8 and, hence, contains no outliers .

## Catagorical Features

```
In [19]: categoricalColumns = df.columns[df.dtypes==object].tolist()

for col in categoricalColumns:
    print('Unique values for ' + col)
    print(df[col].unique())
    print('')
```

```
Unique values for school
['GP' 'MS']
```

```
Unique values for sex
['F' 'M']
```

```
Unique values for address
['U' 'R']
```

```
Unique values for famsize
['GT3' 'LE3']
```

```
Unique values for schoolsup
['yes' 'no']
```

```
Unique values for famsup
['no' 'yes']
```

```
Unique values for paid
['no' 'yes']
```

```
Unique values for higher
['yes' 'no']
```

```
Unique values for internet
['no' 'yes']
```

```
Unique values for romantic
['no' 'yes']
```

Each catagorical feature contains only two unique values each, such as 'yes' or 'no'.

It seems like no accidental symbol, such as a full stop (.) is in any of the rows for any of the catagorical columns, so we dont need to use:

```
df['column_name'].str.rstrip(".")
```

We can now consider our dataset 'clean' & ready for visualisation & data modelling.


## Univariate Visualisation

Lets get a histogram of Study time & Travel times.

```
In [20]: ax = df['studytime'].value_counts().plot(kind = 'bar', color = 'green')
ax.set_xticklabels(ax.get_xticklabels(), rotation = 90)
plt.tight_layout()
plt.title('Figure 1: Bar Chart of Study time', fontsize = 15)
plt.show();
```

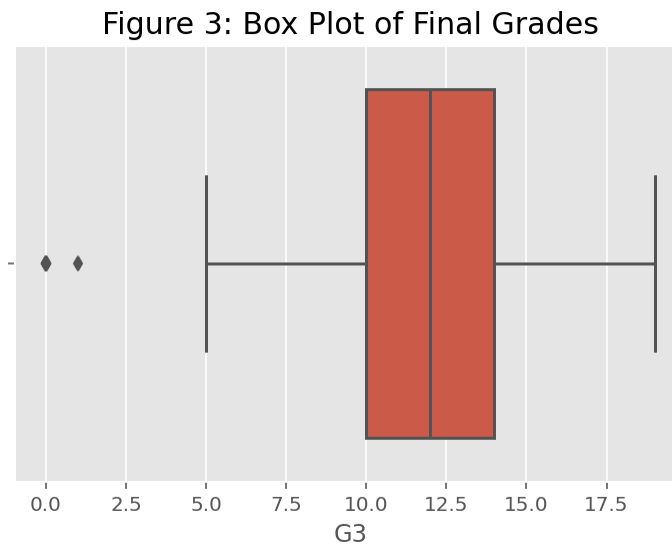


```
In [21]: ax = df['traveltime'].value_counts().plot(kind = 'bar', color = 'red')
ax.set_xticklabels(ax.get_xticklabels(), rotation = 90)
plt.tight_layout()
plt.title('Figure 2: Bar Chart of Travel time', fontsize = 15)
plt.show();
```




Let's display a boxplot and histogram for Final Grades. Figure 3 shows that this variable is left-skewed.

```
In [22]: # get a box plot of final grades
sns.boxplot(df['G3']).set_title('Figure 3: Box Plot of Final Grade
s', fontsize = 15)
plt.show();
```



```
In [23]: # get a histogram of age with kernel density estimate
sns.distplot(df['G3'], kde = True).set_title('Figure 4: Histogram
of Final Grades', fontsize = 15)
plt.show();
```




## Multivariate Visualisation

### Scatterplot of Numerical Features & Final Grades

We will make a scatterplot between absences and final grades using the `absences` from the copy of the unmodified dataframe.

```
In [24]: plt.scatter(absences, df['G3'], alpha = 0.3)
plt.title('Figure 5: Scatterplot of Absences and Final Grades', fo
ntsize = 15)
plt.xlabel('Absences')
plt.ylabel('Final grades')
plt.show();
```

Figure 5: Scatterplot of Absences and Final Grades

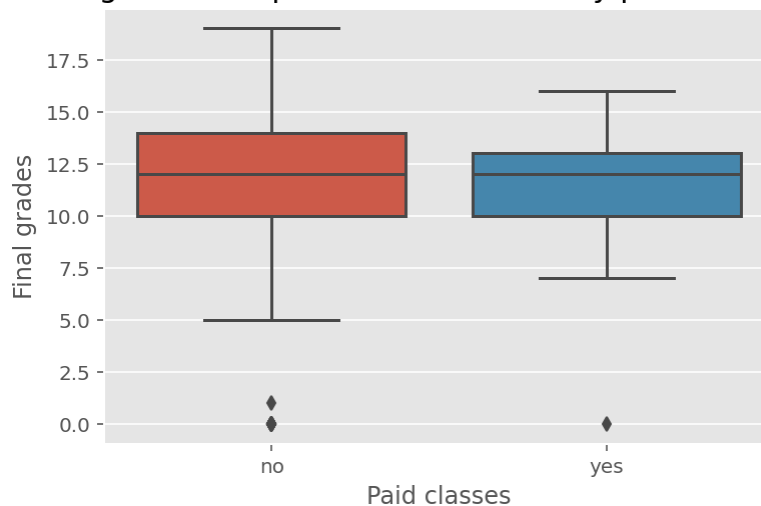


The scatterplot in Figure 5 shows slightest negative correlation between the absences and final grades numeric variables.

## Catagorical attributes by Final Grades

```
In [25]: # Creating a boxplot
sns.boxplot(df['paid'], df['G3']);
plt.title('Figure 6: Boxplot of Final Grades by paid classes', fon
tsize = 15)
plt.xlabel('Paid classes')
plt.ylabel('Final grades')
plt.show();
```

Figure 6: Boxplot of Final Grades by paid classes





The distribution of students taking paid classes and those that aren't does not differ significantly, but have a similar median as seen in Figure 6. The whiskers suggest that students with a final grade of over 16 are more likely to not be involved in any sort of paid classes.

## Statistical Modeling & Performance Evaluation

### Full Model

We begin by fitting a multiple linear regression that predicts final grades using all of the available features. We call this the full model. First let's take a quick peak at the clean data.

In [26]: `df.sample(10)`

Out[26]:

	school	sex	address	famsize	Medu	Fedu	traveltime	studytime	failures	schoolsup
557	MS	F	R	GT3	3	1	1	1	3	no
417	GP	F	U	GT3	3	2	0	2	0	no
318	GP	F	R	GT3	3	2	1	1	0	no
552	MS	M	U	GT3	1	1	0	1	2	no
276	GP	M	U	GT3	2	1	2	0	0	no
310	GP	F	R	GT3	2	1	1	1	0	no
570	MS	M	R	GT3	2	2	1	0	0	no
371	GP	F	U	GT3	2	2	0	0	0	no
553	MS	M	U	LE3	1	0	0	0	0	no
378	GP	M	U	GT3	3	3	0	0	0	no

When constructing the regression formula, we can manually add all the independent features.

In [27]: `dependant_var = 'G3'`  
`independant_var = ' + '.join(df.drop(columns=['G3']).columns)`  
`formula_string = dependant_var + ' ~ ' + independant_var`  
`print('formula_string: ', formula_string)`

```
formula_string: G3 ~ school + sex + address + famsize + Medu + Fe
du + traveltime + studytime + failures + schoolsup + famsup + paid
+ higher + internet + romantic + famrel + freetime + health + abse
nces + Wkalc
```

The formula string above works just fine with the Statsmodels module. The problem, however, is that we cannot do automatic variable selection with this formula. What we need for this purpose is "one-hot-encoding" of categorical features.

In the code chunk below, we first use the `get_dummies()` function in Pandas for one-hot-encoding of categorical features and then we construct a new formula string with the encoded features.

```
In [28]: data_encoded = pd.get_dummies(df, drop_first=True)
data_encoded.head()
```

```
Out[28]:
```

	Medu	Fedu	traveltime	studytime	failures	famrel	freetime	health	absences	G3	Wk
0	4	4	1	1	0	3	2	2	0	11	
1	1	1	0	1	0	4	2	2	0	11	
2	1	1	0	1	0	3	2	2	0	12	
3	4	2	0	2	0	2	1	4	0	14	
4	3	3	0	1	0	3	2	4	0	13	

```
In [29]: formula_string_indep_vars_encoded = ' + '.join(data_encoded.drop(c
columns='G3').columns)
formula_string_encoded = 'G3 ~ ' + formula_string_indep_vars_encoded
print('formula_string_encoded: ', formula_string_encoded)
```

```
formula_string_encoded: G3 ~ Medu + Fedu + traveltime + studytime
+ failures + famrel + freetime + health + absences + Wkalc + schoo
l_MS + sex_M + address_U + famsize_LE3 + schoolsup_yes + famsup_ye
s + paid_yes + higher_yes + internet_yes + romantic_yes
```

Now that we have defined our statistical model formula as a Python string, we fit an OLS (ordinary least squares) model to our encoded data.

```
In [30]: model = sm.formula.ols(formula = formula_string_encoded, data = da
         ta_encoded)
         model_fitted = model.fit()
         print(model_fitted.summary())
```

# OLS Regression Results

```

=====
Dep. Variable:          G3      R-squared:
0.339
Model:                  OLS      Adj. R-squared:
0.318
Method:                  Least Squares      F-statistic:
16.10
Date:                    Sun, 01 Nov 2020      Prob (F-statistic):
2.13e-44
Time:                    22:39:47      Log-Likelihood:
-1547.2
No. Observations:        649      AIC:
3136.
Df Residuals:            628      BIC:
3230.
Df Model:                 20
Covariance Type:         nonrobust
=====

```

```

=====
coef      std err      t      P>|t|      [0.
025      0.975]
-----
Intercept      10.6367      0.697      15.265      0.000      9.
268      12.005
Medu      0.1405      0.128      1.094      0.274      -0.
112      0.393
Fedu      0.2053      0.127      1.614      0.107      -0.
044      0.455
traveltime      0.0649      0.156      0.416      0.677      -0.
241      0.371
studytime      0.4255      0.136      3.135      0.002      0.
159      0.692
failures      -1.2913      0.193      -6.694      0.000      -1.
670      -0.912
famrel      0.1176      0.114      1.035      0.301      -0.
106      0.341
freetime      -0.1304      0.104      -1.255      0.210      -0.
334      0.074
health      -0.1704      0.075      -2.283      0.023      -0.
317      -0.024
absences      -0.3173      0.185      -1.717      0.087      -0.
680      0.046
Wkalc      -0.1549      0.059      -2.630      0.009      -0.
271      -0.039
school_MS      -1.3571      0.255      -5.323      0.000      -1.
858      -0.857
sex_M      -0.5728      0.243      -2.355      0.019      -1.
051      -0.095
address_U      0.2845      0.256      1.110      0.267      -0.
219      0.788
famsize_LE3      0.2858      0.233      1.225      0.221      -0.
172      0.744
schoolsup_yes      -1.4168      0.353      -4.010      0.000      -2.
111      -0.723
famsup_yes      -0.0118      0.224      -0.053      0.958      -0.
451      0.427
paid_yes      -0.5291      0.453      -1.167      0.244      -1.

```

```

420      0.361
higher_yes      1.6709      0.373      4.474      0.000      0.
937      2.404
internet_yes      0.3667      0.268      1.369      0.171      -0.
159      0.893
romantic_yes      -0.3742      0.223      -1.676      0.094      -0.
813      0.064
=====
=====
Omnibus:      108.875      Durbin-Watson:
1.882
Prob(Omnibus):      0.000      Jarque-Bera (JB):
328.833
Skew:      -0.806      Prob(JB):
3.93e-72
Kurtosis:      6.092      Cond. No.
45.4
=====
=====

```

#### Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

The equation of the regression model that includes all of the variables rounded to two decimal places is:

$10.64 + (0.14 \text{ Medu}) + (0.21 \text{ Fedu}) + (0.06 \text{ traveltime}) + (0.43 \text{ studytime}) + (-1.29 \text{ failures}) + (0.12 \text{ famrel}) + (-0.13 \text{ freetime}) + (-0.17 \text{ health}) + (-0.32 \text{ absences}) + (-0.15 \text{ Walc}) + (-1.36 \text{ School\_MS}) + (-0.57 \text{ sex\_M}) + (0.28 \text{ address\_U}) + (0.29 \text{ famsize\_LE3}) + (-1.42 \text{ schoolsup\_yes}) + (-0.01 \text{ famsup\_yes}) + (-0.53 \text{ paid\_yes}) + (1.67 \text{ higher\_yes}) + (0.37 \text{ internet\_yes}) + (-0.37 \text{ romantic\_yes})$

The intercept in this case, in simple terms, refers to the expected final grade if everything else was 0, such as all of the nominal categorical features are all 'no' (eg: famsup\_yes = 0).

Overall, a student is expected to receive 10.6367 as their final grade if their mother and father has no education, home to school travel time is less than 15 minutes, study time is less than 2 hours, they have 0 previous failures, their family relationship is very bad, freetime afterschool is very low, their current health status is very bad, their absences from school is very low, weekly alcohol consumption is very low, they attend Gabriel Pereira school, are females, live in rural areas, have a family size greater than 3, receive no school support, no family support, no extra paid classes, does not want to go into higher education, does not have internet and is not in a romantic relationship.

```
In [31]: residuals_full = pd.DataFrame({'actual': df['G3'],
                                         'predicted': model_fitted.fittedvalue
                                         s,
                                         'residual': model_fitted.resid})
residuals_full.head(10)
```

Out[31]:

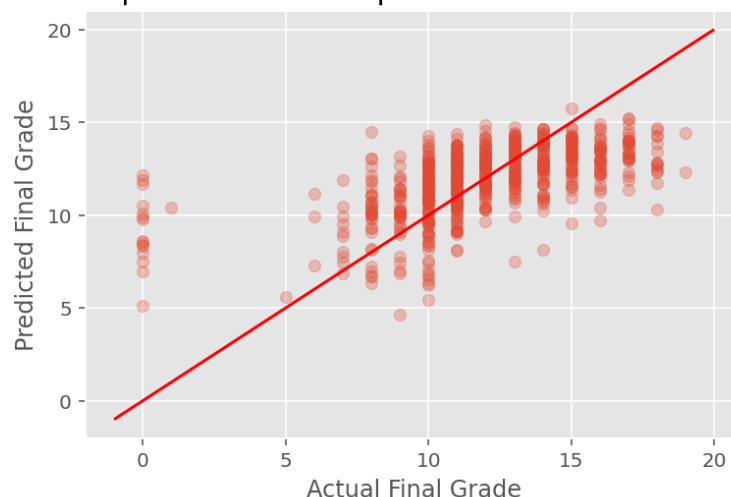
	actual	predicted	residual
0	11	12.800059	-1.800059
1	11	13.587117	-2.587117
2	12	11.885661	0.114339
3	14	13.819668	0.180332
4	13	13.298714	-0.298714
5	13	13.506113	-0.506113
6	13	13.409647	-0.409647
7	13	13.389911	-0.389911
8	17	14.140013	2.859987
9	13	13.309535	-0.309535

Let's plot actual final grade values vs. predicted values.

```
In [32]: def plot_line(axis, slope, intercept, **kargs):
           xmin, xmax = axis.get_xlim()
           plt.plot([xmin, xmax], [xmin*slope+intercept, xmax*slope+intercept], **kargs)

plt.scatter(residuals_full['actual'], residuals_full['predicted'],
            alpha=0.3);
plot_line(axis=plt.gca(), slope=1, intercept=0, c="red");
plt.xlabel('Actual Final Grade');
plt.ylabel('Predicted Final Grade');
plt.title('Figure 7: Scatter plot of actual vs. predicted Final Grade for the Full Model',
          fontsize=15);
plt.show();
```

Figure 7: Scatter plot of actual vs. predicted Final Grade for the Full Model



From Figure 7, we observe that the model never produces a prediction above 16 even though the highest final grade in the dataset is 19.

We will now check the diagnostics for the full model.

## Full Model Diagnostic Checks


We would like to check whether there are indications of violations of the regression assumptions, which are

1. linearity of the relationship between target variable and the independent variables
2. constant variance of the errors
3. normality of the residual distribution
4. statistical independence of the residuals

Let's first get a scatter plot of residuals (as a function of predicted final grades).

```
In [33]: plt.scatter(residuals_full['predicted'], residuals_full['residual'], alpha=0.3);  
plt.xlabel('Predicted Final Grades');  
plt.ylabel('Residuals')  
plt.title('Figure 8(a): Scatterplot of residuals vs. predicted final grades for Full Model', fontsize=15)  
plt.show();
```


Figure 8(a): Scatterplot of residuals vs. predicted final grades for Full Model



Let's now plot actual age vs. residuals.

```
In [34]: plt.scatter(residuals_full['actual'], residuals_full['residual'],
alpha=0.3);
plt.xlabel('Actual Final Grades');
plt.ylabel('Residuals')
plt.title('Figure 8(b): Scatterplot of residuals vs. actual final
grades for Full Model', fontsize=15)
plt.show();
```

Figure 8(b): Scatterplot of residuals vs. actual final grades for Full Model




We notice that the model overestimates lower final grades. In particular, for those with a final grade less than 5, the model predicts much higher final grades.

Let's overlay the histograms of actual vs. predicted final grades on the same plot.

```
In [35]: plt.hist(residuals_full['actual'], label='actual', bins=20, alpha=
0.7);
plt.hist(residuals_full['predicted'], label='predicted', bins=20,
alpha=0.7);
plt.xlabel('Final Grades');
plt.ylabel('Frequency');
plt.title('Figure 9: Histograms of actual final grades vs. predict
ed final grades for Full Model', fontsize=15);
plt.legend()
plt.show();
```

Figure 9: Histograms of actual final grades vs. predicted final grades for Full Model






We notice that their distributions are quite different. In particular, the model's predictions are highly clustered around mid-13's.

Let's now have look at the histogram of the residuals.

```
In [36]: plt.hist(residuals_full['residual'], bins = 20);  
plt.xlabel('Residual');  
plt.ylabel('Frequency');  
plt.title('Figure 10: Histogram of residuals for Full Model', font  
size=15);  
plt.show();
```



From Figure 10, the histogram of residuals looks somewhat symmetric, though slightly left-skewed. Nonetheless, it seems the normality assumption of linear regression is not significantly violated in this particular case.

## Backwards Feature Selection

We now perform backwards feature selection using p-values. It appears Statsmodels does not have any canned code for automatic feature selection, so we wrote one ourselves.

```
In [37]: patsy_description = patsy.ModelDesc.from_formula(formula_string_encoded)
linreg_fit = model_fitted
p_val_cutoff = 0.05

print('\nPerforming backwards feature selection using p-values:')

while True:
    pval_series = linreg_fit.pvalues.drop(labels='Intercept')
    pval_series = pval_series.sort_values(ascending=False)
    term = pval_series.index[0]
    pval = pval_series[0]
    if (pval < p_val_cutoff):
        break
    term_components = term.split(':')
    print(f'\nRemoving term "{term}" with p-value {pval:.4f}')
    if (len(term_components) == 1):
        patsy_description.rhs_termlist.remove(patsy.Term([patsy.EvalFactor(term_components[0])]))
    else:
        patsy_description.rhs_termlist.remove(patsy.Term([patsy.EvalFactor(term_components[0]),
                                                                    patsy.EvalFactor(term_components[1])]))

    linreg_fit = smf.ols(formula=patsy_description, data=data_encoded).fit()

model_reduced_fitted = smf.ols(formula = patsy_description, data = data_encoded).fit()

print("\n***")
print(model_reduced_fitted.summary())
print("****")
print(f"Regression number of terms: {len(model_reduced_fitted.model.exog_names)}")
print(f"Regression F-distribution p-value: {model_reduced_fitted.f_pvalue:.4f}")
print(f"Regression R-squared: {model_reduced_fitted.rsquared:.4f}")
print(f"Regression Adjusted R-squared: {model_reduced_fitted.rsquared_adj:.4f}")
```

Performing backwards feature selection using p-values:

Removing term "famsup\_yes" with p-value 0.958

Removing term "traveltime" with p-value 0.6775

Removing term "famrel" with p-value 0.2995

Removing term "address\_U" with p-value 0.3248

Removing term "Medu" with p-value 0.2688

Removing term "paid\_yes" with p-value 0.2576

Removing term "freetime" with p-value 0.2651

Removing term "famsize\_LE3" with p-value 0.1636

Removing term "internet\_yes" with p-value 0.1073

Removing term "absences" with p-value 0.1222

Removing term "romantic\_yes" with p-value 0.07934

Removing term "sex\_M" with p-value 0.05772

\*\*\*

#### OLS Regression Results

```

=====
=====
Dep. Variable:          G3    R-squared:
0.318
Model:                OLS    Adj. R-squared:
0.310
Method:              Least Squares    F-statistic:
37.34
Date:                Sun, 01 Nov 2020    Prob (F-statistic):
1.06e-48
Time:                22:39:48    Log-Likelihood:
-1557.2
No. Observations:    649    AIC:
3132.
Df Residuals:        640    BIC:
3173.
Df Model:              8
Covariance Type:      nonrobust
=====
=====

```

		coef	std err	t	P> t	[0.
025	0.975]					
Intercept		10.9878	0.481	22.825	0.000	10.
043	11.933					
Fedu		0.2755	0.100	2.749	0.006	0.
079	0.472					
studytime		0.4806	0.133	3.605	0.000	0.
219	0.742					
failures		-1.4381	0.190	-7.574	0.000	-1.
811	-1.065					

```

health      -0.1912      0.074      -2.597      0.010      -0.
336      -0.047
Wkalc      -0.2162      0.055      -3.953      0.000      -0.
324      -0.109
school_MS   -1.4241      0.231      -6.170      0.000      -1.
877      -0.971
schoolsup_yes -1.3615      0.349      -3.899      0.000      -2.
047      -0.676
higher_yes   1.8616      0.370      5.028      0.000      1.
135      2.589

```

```

=====
=====
Omnibus:      103.221   Durbin-Watson:
1.918
Prob(Omnibus):      0.000   Jarque-Bera (JB):
301.889
Skew:      -0.775   Prob(JB):
2.79e-66
Kurtosis:      5.960   Cond. No.
23.7
=====
=====

```

#### Warnings:

```
[1] Standard Errors assume that the covariance matrix of the error
s is correctly specified.
```

```
***
```

```
Regression number of terms: 9
```

```
Regression F-distribution p-value: 0.0000
```

```
Regression R-squared: 0.3182
```

```
Regression Adjusted R-squared: 0.3097
```

Similar to what we did for the full model, let's define a new data frame for actual final grade vs. predicted final grade and the residuals for the reduced model.


```
In [38]: residuals_reduced = pd.DataFrame({'actual': data_encoded['G3'],
                                           'predicted': model_fitted.fittedvalues,
                                           'residual': model_fitted.resid})
residuals_reduced.head(10)
```

Out[38]:

	actual	predicted	residual
0	11	12.800059	-1.800059
1	11	13.587117	-2.587117
2	12	11.885661	0.114339
3	14	13.819668	0.180332
4	13	13.298714	-0.298714
5	13	13.506113	-0.506113
6	13	13.409647	-0.409647
7	13	13.389911	-0.389911
8	17	14.140013	2.859987
9	13	13.309535	-0.309535

```
In [39]: plt.scatter(residuals_reduced['actual'], residuals_reduced['predicted'], alpha=0.3);
plot_line(axis=plt.gca(), slope=1, intercept=0, c="red");
plt.xlabel('Final Grade');
plt.ylabel('Predicted Final Grade');
plt.title('Figure 11: Scatter plot of actual vs. predicted Final Grade for Reduced Model', fontsize=15);
plt.show();
```

Figure 11: Scatter plot of actual vs. predicted Final Grade for Reduced Model




This model returns an Adjusted R-squared of 0.3097, meaning the reduced model still explains about 31% of the variance, but with 6 less variables. Looking at the p-values, they are all significant at the 5% level, as expected. From Figure 11, we still have the same issues with our model. That is, the model overestimates higher grades and underestimates lower grades. We will now perform the diagnostic checks on this reduced model.

## Reduced Model


```
In [40]: plt.scatter(residuals_reduced['predicted'], residuals_reduced['residual'], alpha=0.3);  
plt.xlabel('Predicted Final Grade');  
plt.ylabel('Residuals')  
plt.title('Figure 12: Scatter plot of residuals vs. predicted Final Grade for Reduced Model', fontsize=15)  
plt.show();
```

Figure 12: Scatter plot of residuals vs. predicted Final Grade for Reduced Model



```
In [41]: plt.hist(residuals_reduced['residual'], bins = 20);  
plt.xlabel('Residual');  
plt.ylabel('Frequency')  
plt.title('Figure 13: Histogram of residuals for Reduced Model', fontsize = 15)  
plt.show();
```

Figure 13: Histogram of residuals for Reduced Model



## Summary & Conclusions

Using our independent variables, we were able to get a full model with an Adjusted R-squared value of about 31%. After backwards variable selection with a p-value cutoff value of 0.05, we were able to maintain the same performance but with 12 less variables. Our final model has 9 variables all together with a model p-value of 0.

The final multiple linear regression model has an Adjusted R-squared value of about 31%, which is significantly low. So, it appears that the variables we used are not enough for accurately estimating the final grade of students from a dataset of student achievement from two Portuguese schools. Next time we should add some more interaction terms and maybe some other higher order terms to see if this would result in some improvement for the Adjusted R-squared value. Nonetheless, it might be the case that nonlinear models such as a neural network might be more appropriate for the task at hand rather than a linear regression model. Our regression model appears to predict final grade correctly within 10-13 marks in general, though this is clearly a huge margin of error for the model to be useful for any practical purposes.

Furthermore, our model has some quite significant problems. More specifically, our model consistently underestimates lower final grades and overestimates higher final grades. In particular, for those who received a final grade below 8, the model predicts much lower grades. Also, for those who received a final grade above 15, the model predicts significantly higher final grades.

## References

- P. Cortez and A. Silva. Using Data Mining to Predict Secondary School Student Performance. In A. Brito and J. Teixeira Eds., Proceedings of 5th FUTURE BUSINESS TECHNOLOGY Conference (FUBUTEC 2008) pp. 5-12, Porto, Portugal, April, 2008, EUROSIS, ISBN 978-9077381-39-7. Available at <http://archive.ics.uci.edu/ml/datasets/Student+Performance> (<http://archive.ics.uci.edu/ml/datasets/Student+Performance>) [Accessed 2020-20-07]
- Regression Case Study : Predicting Age in Census Data. Available at <https://www.featureranking.com/tutorials/statistics-tutorials/regression-case-study/> (<https://www.featureranking.com/tutorials/statistics-tutorials/regression-case-study/>) [Accessed 2020-20-07]
- Data Preparation for Statistical Modeling and Machine Learning. Available at <https://www.featureranking.com/tutorials/machine-learning-tutorials/data-preparation-for-machine-learning/> (<https://www.featureranking.com/tutorials/machine-learning-tutorials/data-preparation-for-machine-learning/>) [Accessed 2020-20-07]

In [ ]: